

Answer the questions below. For the questions involving both programming and analysis, please submit your code through Udacity and your analysis through T-square.

1. In a list of distinct numbers a_1, \dots, a_n , we say that the elements a_i and a_j are *inverted* if $i < j$ but $a_i > a_j$. Suppose that all orderings of a_1, \dots, a_n are equally probable under some probability distribution (In other words we shuffled a set of numbers perfectly to obtain the order $a_1 \dots a_n$.) What is the expected number of inversions?
2. Consider the following greedy algorithm for finding a matching.
 - Initialize S to an empty set of edges
 - While there is an edge where *both* vertices are unmatched by S
 - Add the above edge to S
 - Return S
 - a. What is the running time of this algorithm?
 - b. Give an example of a graph where the algorithm **is not guaranteed to** find the maximum matching.
 - c. Show that the matching found by the algorithm always has at least half as many edges as a maximum matching.
 - d. (Bonus) Now consider a similar algorithm for finding a maximum weight matching in an edge-weighted graph: Greedily add the heaviest edge possible to the current matching; stop when no further edge can be added. Show that this algorithm finds a matching whose weight is at least half the optimum.
3. Consider the following factor 2 approximation algorithm for the minimum cardinality vertex cover problem. Find a depth first search tree in the given graph, G , and output the set, say S , of all the non-leaf vertices of this tree. Show that S is indeed a vertex cover for G and $|S| \leq 2 \cdot OPT$.

Hint: Use the tree to construct a matching that includes all the vertices in S .

4. The following is known as the maximum acyclic subgraph problem. Given a directed graph $G = (V, E)$, pick a maximum cardinality set of edges from E so that the resulting subgraph is acyclic. Implement an algorithm that gives a $1/2$ factor approximation here.

<https://www.udacity.com/course/viewer#!/c-ud557/l-1209378918/m-3519618629>

Hint: Rather than thinking about picking $1/2$ the optimum, think about picking $1/2$ of the edges.

5. In the lesson, we gave a randomized algorithm for finding a minimum cut set in graph. Now, we consider the problem of finding a *maximum* cut set (all weights on edges are all one).
- Consider an algorithm that partitions the vertices into two sets A and B by placing each vertex uniformly and independently at random into one of the two sets. What is the expected value of $C(A, B)$, i.e. the number of edges crossing the cut?
 - Use the method of conditional expectation to derandomize the above algorithm so that it always achieves this expected value. Implement your algorithm here.

<https://www.udacity.com/course/viewer#!/c-ud557/l-1209378918/m-3481888754>

- Explain why your algorithm is correct. Use the notation $x_k = A$ to describe the event that vertex k is assigned to the set A in the partition. Use Y to denote the number of cut edges, and let $P : \{1, \dots, n\} \rightarrow \{A, B\}$ represent the partition returned by the derandomized algorithm.